

IEOR142_Project

May 8, 2023

1 INDENG 142 - Final Project

Grace Jung, Vijay Baliga, Nicholas Solis, Robinson Kao, David Ko

```
[1]: import random
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
import statsmodels.formula.api as smf
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import OneHotEncoder
```

import dataset excel file as a dataframe

```
[2]: df = pd.read_csv(r"Data Science Jobs Salaries.csv")
```

```
[3]: df
```

```
[3]:
```

	work_year	experience_level	employment_type	job_title	\
0	2021e	EN	FT	Data Science Consultant	
1	2020	SE	FT	Data Scientist	
2	2021e	EX	FT	Head of Data Science	
3	2021e	EX	FT	Head of Data	
4	2021e	EN	FT	Machine Learning Engineer	
..	
240	2020	SE	FT	Data Scientist	
241	2021e	MI	FT	Principal Data Scientist	
242	2020	EN	FT	Data Scientist	
243	2020	EN	CT	Business Data Analyst	
244	2021e	SE	FT	Data Science Manager	

	salary	salary_currency	salary_in_usd	employee_residence	remote_ratio	\
0	54000	EUR	64369	DE	50	
1	60000	EUR	68428	GR	100	
2	85000	USD	85000	RU	0	
3	230000	USD	230000	RU	50	
4	125000	USD	125000	US	100	

```

..      ...      ...      ...      ...      ...
240    412000      USD      412000      US      100
241    151000      USD      151000      US      100
242    105000      USD      105000      US      100
243    100000      USD      100000      US      100
244    7000000      INR      94917      IN      50

```

```

      company_location company_size
0                DE             L
1                US             L
2                RU             M
3                RU             L
4                US             S
..      ...      ...
240        US             L
241        US             L
242        US             S
243        US             L
244        IN             L

```

```
[245 rows x 11 columns]
```

2 Data Preprocessing

2.1 drop unnecessary columns

```
[4]: df.drop(columns=['salary_currency', 'salary'], inplace=True)
```

```
[5]: df = df.dropna()
```

2.2 one-hot encode categorical variables

```
[6]: df = pd.get_dummies(df)
```

```
[7]: for col in df.columns:
      df.rename(columns={col : '_' + col.split('_')[1]}, inplace=True)
df.columns
```

```
[7]: Index(['salary_in_usd', 'remote_ratio', 'work_year_2020', 'work_year_2021e',
          'experience_level_EN', 'experience_level_EX', 'experience_level_MI',
          'experience_level_SE', 'employment_type_CT', 'employment_type_FL',
          ...,
          'company_location_RU', 'company_location_SG', 'company_location_SI',
          'company_location_TR', 'company_location_UA', 'company_location_US',
          'company_location_VN', 'company_size_L', 'company_size_M',
          'company_size_S'],
          dtype=object)
```

```
dtype='object', length=144)
```

2.3 Create n-1 columns for categorical variables to prevent the dummy variable trap. Drop 1 of each.

```
[8]: df.drop(columns=['work_year_2020', 'experience_level_EN', 'employment_type_CT', 'job_title_3D_Computer_Vision_Researcher', 'employee_residence_AE', 'company_location_AE', 'company_size_L'], inplace=True)
```

2.4 Process DataFrame into more consistent, clear categorical variables.

```
[9]: df
```

```
[9]:
```

	salary_in_usd	remote_ratio	work_year_2021e	experience_level_EX	\
0	64369	50	1	0	
1	68428	100	0	0	
2	85000	0	1	1	
3	230000	50	1	1	
4	125000	100	1	0	
..	
240	412000	100	0	0	
241	151000	100	1	0	
242	105000	100	0	0	
243	100000	100	0	0	
244	94917	50	1	0	

	experience_level_MI	experience_level_SE	employment_type_FL	\
0	0	0	0	
1	0	1	0	
2	0	0	0	
3	0	0	0	
4	0	0	0	
..	
240	0	1	0	
241	1	0	0	
242	0	0	0	
243	0	0	0	
244	0	1	0	

	employment_type_FT	employment_type_PT	job_title_AI_Scientist	...	\
0	1	0	0	...	
1	1	0	0	...	
2	1	0	0	...	
3	1	0	0	...	
4	1	0	0	...	
..	

240	1	0	0	...
241	1	0	0	...
242	1	0	0	...
243	0	0	0	...
244	1	0	0	...

	company_location_PT	company_location_RU	company_location_SG	\
0	0	0	0	
1	0	0	0	
2	0	1	0	
3	0	1	0	
4	0	0	0	
..	
240	0	0	0	
241	0	0	0	
242	0	0	0	
243	0	0	0	
244	0	0	0	

	company_location_SI	company_location_TR	company_location-UA	\
0	0	0	0	
1	0	0	0	
2	0	0	0	
3	0	0	0	
4	0	0	0	
..	
240	0	0	0	
241	0	0	0	
242	0	0	0	
243	0	0	0	
244	0	0	0	

	company_location_US	company_location_VN	company_size_M	company_size_S
0	0	0	0	0
1	1	0	0	0
2	0	0	1	0
3	0	0	0	0
4	1	0	0	1
..
240	1	0	0	0
241	1	0	0	0
242	1	0	0	1
243	1	0	0	0
244	0	0	0	0

[245 rows x 137 columns]

```
[10]: # Unique values of each column
print('COLUMN NULL COUNT')
column_types = {}
for col in df.columns:
    # print(i)
    # print(df[i].unique())
    column_types[col] = type(df[col][0])
    null_count = sum(df[col].isna())
    print(col + ': ', null_count)

    # print('-----')

print('\nCOLUMN DATA TYPE')
for col, type in column_types.items():
    print(col + ': ', type)
```

```
COLUMN NULL COUNT
salary_in_usd: 0
remote_ratio: 0
work_year_2021e: 0
experience_level_EX: 0
experience_level_MI: 0
experience_level_SE: 0
employment_type_FL: 0
employment_type_FT: 0
employment_type_PT: 0
job_title_AI_Scientist: 0
job_title_Applied_Data_Scientist: 0
job_title_Applied_Machine_Learning_Scientist: 0
job_title_BI_Data_Analyst: 0
job_title_Big_Data_Architect: 0
job_title_Big_Data_Engineer: 0
job_title_Business_Data_Analyst: 0
job_title_Cloud_Data_Engineer: 0
job_title_Computer_Vision_Engineer: 0
job_title_Computer_Vision_Software_Engineer: 0
job_title_Data_Analyst: 0
job_title_Data_Analytics_Engineer: 0
job_title_Data_Analytics_Manager: 0
job_title_Data_Architect: 0
job_title_Data_Engineer: 0
job_title_Data_Engineering_Manager: 0
job_title_Data_Science_Consultant: 0
job_title_Data_Science_Engineer: 0
job_title_Data_Science_Manager: 0
job_title_Data_Scientist: 0
job_title_Data_Specialist: 0
```

job_title_Director_of_Data_Engineering: 0
job_title_Director_of_Data_Science: 0
job_title_Finance_Data_Analyst: 0
job_title_Financial_Data_Analyst: 0
job_title_Head_of_Data: 0
job_title_Head_of_Data_Science: 0
job_title_Lead_Data_Analyst: 0
job_title_Lead_Data_Engineer: 0
job_title_Lead_Data_Scientist: 0
job_title_ML_Engineer: 0
job_title_Machine_Learning_Engineer: 0
job_title_Machine_Learning_Infrastructure_Engineer: 0
job_title_Machine_Learning_Scientist: 0
job_title_Manager_Data_Science: 0
job_title_Marketing_Data_Analyst: 0
job_title_Principal_Data_Analyst: 0
job_title_Principal_Data_Engineer: 0
job_title_Principal_Data_Scientist: 0
job_title_Product_Data_Analyst: 0
job_title_Research_Scientist: 0
job_title_Staff_Data_Scientist: 0
employee_residence_AT: 0
employee_residence_BE: 0
employee_residence_BG: 0
employee_residence_BR: 0
employee_residence_CA: 0
employee_residence_CL: 0
employee_residence_CN: 0
employee_residence_CO: 0
employee_residence_DE: 0
employee_residence_DK: 0
employee_residence_ES: 0
employee_residence_FR: 0
employee_residence_GB: 0
employee_residence_GR: 0
employee_residence_HK: 0
employee_residence_HR: 0
employee_residence_HU: 0
employee_residence_IN: 0
employee_residence_IR: 0
employee_residence_IT: 0
employee_residence_JE: 0
employee_residence_JP: 0
employee_residence_KE: 0
employee_residence_LU: 0
employee_residence_MD: 0
employee_residence_MT: 0
employee_residence_MX: 0

employee_residence_NG: 0
employee_residence_NL: 0
employee_residence_NZ: 0
employee_residence_PH: 0
employee_residence_PK: 0
employee_residence_PL: 0
employee_residence_PR: 0
employee_residence_PT: 0
employee_residence_RO: 0
employee_residence_RS: 0
employee_residence_RU: 0
employee_residence_SG: 0
employee_residence_SI: 0
employee_residence_TR: 0
employee_residence_UA: 0
employee_residence_US: 0
employee_residence_VN: 0
company_location_AS: 0
company_location_AT: 0
company_location_BE: 0
company_location_BR: 0
company_location_CA: 0
company_location_CH: 0
company_location_CL: 0
company_location_CN: 0
company_location_CO: 0
company_location_DE: 0
company_location_DK: 0
company_location_ES: 0
company_location_FR: 0
company_location_GB: 0
company_location_GR: 0
company_location_HR: 0
company_location_HU: 0
company_location_IL: 0
company_location_IN: 0
company_location_IR: 0
company_location_IT: 0
company_location_JP: 0
company_location_KE: 0
company_location_LU: 0
company_location_MD: 0
company_location_MT: 0
company_location_MX: 0
company_location_NG: 0
company_location_NL: 0
company_location_NZ: 0
company_location_PK: 0

company_location_PL: 0
company_location_PT: 0
company_location_RU: 0
company_location_SG: 0
company_location_SI: 0
company_location_TR: 0
company_location_UA: 0
company_location_US: 0
company_location_VN: 0
company_size_M: 0
company_size_S: 0

COLUMN DATA TYPE

salary_in_usd: <class 'numpy.int64'>
remote_ratio: <class 'numpy.int64'>
work_year_2021e: <class 'numpy.uint8'>
experience_level_EX: <class 'numpy.uint8'>
experience_level_MI: <class 'numpy.uint8'>
experience_level_SE: <class 'numpy.uint8'>
employment_type_FL: <class 'numpy.uint8'>
employment_type_FT: <class 'numpy.uint8'>
employment_type_PT: <class 'numpy.uint8'>
job_title_AI_Scientist: <class 'numpy.uint8'>
job_title_Applied_Data_Scientist: <class 'numpy.uint8'>
job_title_Applied_Machine_Learning_Scientist: <class 'numpy.uint8'>
job_title_BI_Data_Analyst: <class 'numpy.uint8'>
job_title_Big_Data_Architect: <class 'numpy.uint8'>
job_title_Big_Data_Engineer: <class 'numpy.uint8'>
job_title_Business_Data_Analyst: <class 'numpy.uint8'>
job_title_Cloud_Data_Engineer: <class 'numpy.uint8'>
job_title_Computer_Vision_Engineer: <class 'numpy.uint8'>
job_title_Computer_Vision_Software_Engineer: <class 'numpy.uint8'>
job_title_Data_Analyst: <class 'numpy.uint8'>
job_title_Data_Analytics_Engineer: <class 'numpy.uint8'>
job_title_Data_Analytics_Manager: <class 'numpy.uint8'>
job_title_Data_Architect: <class 'numpy.uint8'>
job_title_Data_Engineer: <class 'numpy.uint8'>
job_title_Data_Engineering_Manager: <class 'numpy.uint8'>
job_title_Data_Science_Consultant: <class 'numpy.uint8'>
job_title_Data_Science_Engineer: <class 'numpy.uint8'>
job_title_Data_Science_Manager: <class 'numpy.uint8'>
job_title_Data_Scientist: <class 'numpy.uint8'>
job_title_Data_Specialist: <class 'numpy.uint8'>
job_title_Director_of_Data_Engineering: <class 'numpy.uint8'>
job_title_Director_of_Data_Science: <class 'numpy.uint8'>
job_title_Finance_Data_Analyst: <class 'numpy.uint8'>
job_title_Financial_Data_Analyst: <class 'numpy.uint8'>
job_title_Head_of_Data: <class 'numpy.uint8'>

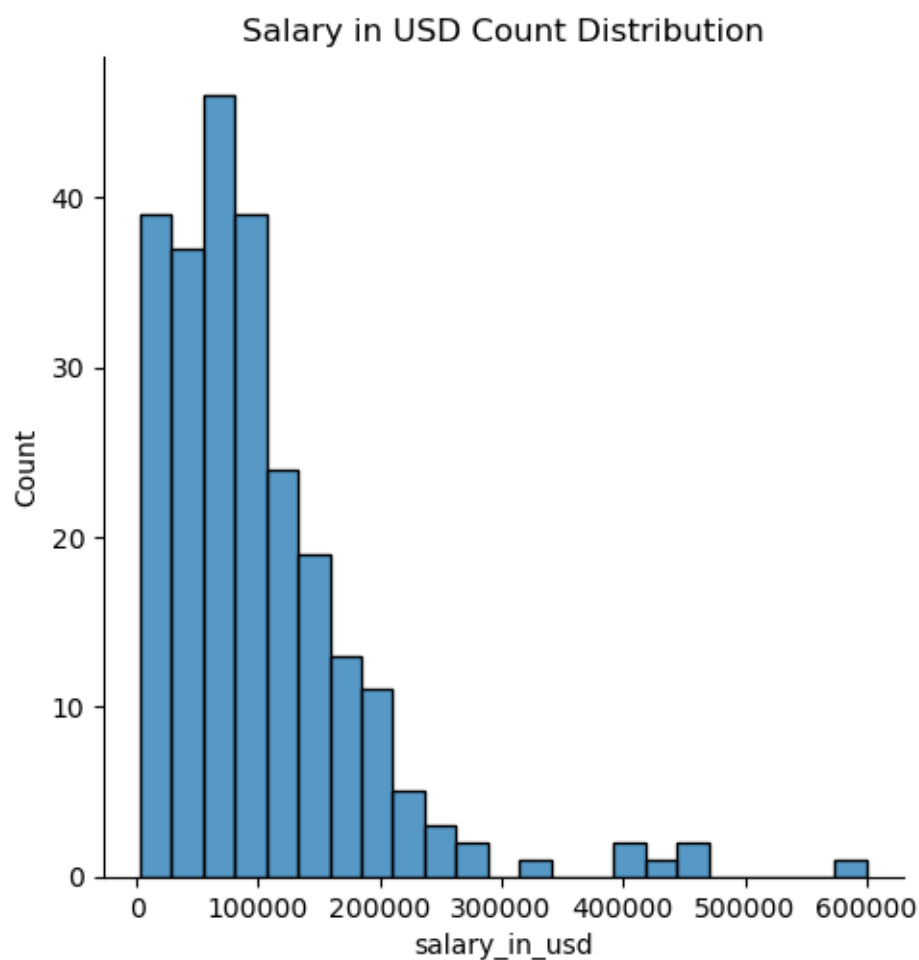
job_title_Head_of_Data_Science: <class 'numpy.uint8'>
job_title_Lead_Data_Analyst: <class 'numpy.uint8'>
job_title_Lead_Data_Engineer: <class 'numpy.uint8'>
job_title_Lead_Data_Scientist: <class 'numpy.uint8'>
job_title_ML_Engineer: <class 'numpy.uint8'>
job_title_Machine_Learning_Engineer: <class 'numpy.uint8'>
job_title_Machine_Learning_Infrastructure_Engineer: <class 'numpy.uint8'>
job_title_Machine_Learning_Scientist: <class 'numpy.uint8'>
job_title_Manager_Data_Science: <class 'numpy.uint8'>
job_title_Marketing_Data_Analyst: <class 'numpy.uint8'>
job_title_Principal_Data_Analyst: <class 'numpy.uint8'>
job_title_Principal_Data_Engineer: <class 'numpy.uint8'>
job_title_Principal_Data_Scientist: <class 'numpy.uint8'>
job_title_Product_Data_Analyst: <class 'numpy.uint8'>
job_title_Research_Scientist: <class 'numpy.uint8'>
job_title_Staff_Data_Scientist: <class 'numpy.uint8'>
employee_residence_AT: <class 'numpy.uint8'>
employee_residence_BE: <class 'numpy.uint8'>
employee_residence_BG: <class 'numpy.uint8'>
employee_residence_BR: <class 'numpy.uint8'>
employee_residence_CA: <class 'numpy.uint8'>
employee_residence_CL: <class 'numpy.uint8'>
employee_residence_CN: <class 'numpy.uint8'>
employee_residence_CO: <class 'numpy.uint8'>
employee_residence_DE: <class 'numpy.uint8'>
employee_residence_DK: <class 'numpy.uint8'>
employee_residence_ES: <class 'numpy.uint8'>
employee_residence_FR: <class 'numpy.uint8'>
employee_residence_GB: <class 'numpy.uint8'>
employee_residence_GR: <class 'numpy.uint8'>
employee_residence_HK: <class 'numpy.uint8'>
employee_residence_HR: <class 'numpy.uint8'>
employee_residence_HU: <class 'numpy.uint8'>
employee_residence_IN: <class 'numpy.uint8'>
employee_residence_IR: <class 'numpy.uint8'>
employee_residence_IT: <class 'numpy.uint8'>
employee_residence_JE: <class 'numpy.uint8'>
employee_residence_JP: <class 'numpy.uint8'>
employee_residence_KE: <class 'numpy.uint8'>
employee_residence_LU: <class 'numpy.uint8'>
employee_residence_MD: <class 'numpy.uint8'>
employee_residence_MT: <class 'numpy.uint8'>
employee_residence_MX: <class 'numpy.uint8'>
employee_residence_NG: <class 'numpy.uint8'>
employee_residence_NL: <class 'numpy.uint8'>
employee_residence_NZ: <class 'numpy.uint8'>
employee_residence_PH: <class 'numpy.uint8'>
employee_residence_PK: <class 'numpy.uint8'>

employee_residence_PL: <class 'numpy.uint8'>
employee_residence_PR: <class 'numpy.uint8'>
employee_residence_PT: <class 'numpy.uint8'>
employee_residence_R0: <class 'numpy.uint8'>
employee_residence_RS: <class 'numpy.uint8'>
employee_residence_RU: <class 'numpy.uint8'>
employee_residence_SG: <class 'numpy.uint8'>
employee_residence_SI: <class 'numpy.uint8'>
employee_residence_TR: <class 'numpy.uint8'>
employee_residence_UA: <class 'numpy.uint8'>
employee_residence_US: <class 'numpy.uint8'>
employee_residence_VN: <class 'numpy.uint8'>
company_location_AS: <class 'numpy.uint8'>
company_location_AT: <class 'numpy.uint8'>
company_location_BE: <class 'numpy.uint8'>
company_location_BR: <class 'numpy.uint8'>
company_location_CA: <class 'numpy.uint8'>
company_location_CH: <class 'numpy.uint8'>
company_location_CL: <class 'numpy.uint8'>
company_location_CN: <class 'numpy.uint8'>
company_location_CO: <class 'numpy.uint8'>
company_location_DE: <class 'numpy.uint8'>
company_location_DK: <class 'numpy.uint8'>
company_location_ES: <class 'numpy.uint8'>
company_location_FR: <class 'numpy.uint8'>
company_location_GB: <class 'numpy.uint8'>
company_location_GR: <class 'numpy.uint8'>
company_location_HR: <class 'numpy.uint8'>
company_location_HU: <class 'numpy.uint8'>
company_location_IL: <class 'numpy.uint8'>
company_location_IN: <class 'numpy.uint8'>
company_location_IR: <class 'numpy.uint8'>
company_location_IT: <class 'numpy.uint8'>
company_location_JP: <class 'numpy.uint8'>
company_location_KE: <class 'numpy.uint8'>
company_location_LU: <class 'numpy.uint8'>
company_location_MD: <class 'numpy.uint8'>
company_location_MT: <class 'numpy.uint8'>
company_location_MX: <class 'numpy.uint8'>
company_location_NG: <class 'numpy.uint8'>
company_location_NL: <class 'numpy.uint8'>
company_location_NZ: <class 'numpy.uint8'>
company_location_PK: <class 'numpy.uint8'>
company_location_PL: <class 'numpy.uint8'>
company_location_PT: <class 'numpy.uint8'>
company_location_RU: <class 'numpy.uint8'>
company_location_SG: <class 'numpy.uint8'>
company_location_SI: <class 'numpy.uint8'>

```
company_location_TR: <class 'numpy.uint8'>
company_location-UA: <class 'numpy.uint8'>
company_location_US: <class 'numpy.uint8'>
company_location_VN: <class 'numpy.uint8'>
company_size_M: <class 'numpy.uint8'>
company_size_S: <class 'numpy.uint8'>
```


2.5 Split dataset into train and test sets.

```
[11]: sns.displot(df, x='salary_in_usd')
plt.title('Salary in USD Count Distribution')
plt.show();
```



```
[12]: # df['log_salary_in_usd'] = np.log(df['salary_in_usd'].values)
# sns.displot(df, x='log_salary_in_usd')
# plt.title('Log Salary in USD Count Distribution');
```

```
[13]: df['cbrt_salary_in_usd'] = np.array(df['salary_in_usd'].values)**(1/3) # Cube Root
sns.displot(data=df, x='cbrt_salary_in_usd')
plt.title('Cube Root Salary in USD Distribution')
plt.show();
```



```
[14]: # train test split
df_train, df_test = train_test_split(df, test_size=0.3, random_state=88)
df_train_y_actual = df_train['cbrt_salary_in_usd']
df_test_y_actual = df_test['cbrt_salary_in_usd']
df_train.shape, df_test.shape
```

```
[14]: ((171, 138), (74, 138))
```

```
[15]: # target_column = 'log_salary_in_usd'
target_column = 'cbrt_salary_in_usd'
feature_columns = [col for col in df_train.columns if col != target_column]
```

```

all_features = ' + '.join(feature_columns)
all_features
linreg = smf.ols(formula = target_column + ' ~ ' + all_features,
                  data = df_train).fit()
print(linreg.summary())
linreg.summary()

```

OLS Regression Results

```

=====
Dep. Variable:      cbrr_salary_in_usd      R-squared:      0.981
Model:              OLS                    Adj. R-squared: 0.954
Method:             Least Squares          F-statistic:   36.26
Date:               Mon, 08 May 2023        Prob (F-statistic): 1.09e-37
Time:               00:59:09                Log-Likelihood: -341.37
No. Observations:   171                    AIC:          884.7
Df Residuals:       70                     BIC:          1202.
Df Model:           100
Covariance Type:    nonrobust
=====

```

```

=====
t      P>|t|      [0.025      0.975]      coef      std err
-----
Intercept      21.1044      4.810
4.388      0.000      11.512      30.697
salary_in_usd      9.947e-05      6.4e-06
15.532      0.000      8.67e-05      0.000
remote_ratio      -0.0072      0.010
-0.746      0.458      -0.026      0.012
work_year_2021e      -0.0626      0.707
-0.089      0.930      -1.473      1.348
experience_level_EX      2.7895      2.897
0.963      0.339      -2.989      8.568
experience_level_MI      0.8204      0.860
0.954      0.343      -0.894      2.535
experience_level_SE      2.0696      1.071
1.932      0.057      -0.067      4.206
employment_type_FL      -7.3516      6.302
-1.167      0.247      -19.920      5.217
employment_type_FT      6.9583      4.491
1.549      0.126      -1.999      15.916
employment_type_PT      7.5817      5.592
1.356      0.179      -3.571      18.734
job_title_AI_Scientist      1.6816      2.945
0.571      0.570      -4.193      7.556
job_title_Applied_Data_Scientist      7.3678      5.312

```

1.387	0.170	-3.227	17.963		
job_title_Applied_Machine_Learning_Scientist				-2.9716	3.282
-0.906	0.368	-9.517	3.574		
job_title_BI_Data_Analyst				1.5084	3.069
0.491	0.625	-4.613	7.630		
job_title_Big_Data_Architect				-0.0997	3.425
-0.029	0.977	-6.931	6.732		
job_title_Big_Data_Engineer				-1.1441	2.279
-0.502	0.617	-5.690	3.402		
job_title_Business_Data_Analyst				-5.588e-11	4.69e-11
-1.192	0.237	-1.49e-10	3.76e-11		
job_title_Cloud_Data_Engineer				-3.146e-11	3.04e-11
-1.034	0.305	-9.21e-11	2.92e-11		
job_title_Computer_Vision_Engineer				0.1228	3.424
0.036	0.971	-6.706	6.951		
job_title_Computer_Vision_Software_Engineer				1.0060	2.969
0.339	0.736	-4.916	6.928		
job_title_Data_Analyst				-0.4791	1.190
-0.403	0.688	-2.852	1.894		
job_title_Data_Analytics_Engineer				1.3717	2.883
0.476	0.636	-4.378	7.121		
job_title_Data_Analytics_Manager				2.0740	2.921
0.710	0.480	-3.751	7.900		
job_title_Data_Architect				2.9565	2.843
1.040	0.302	-2.715	8.628		
job_title_Data_Engineer				1.8334	0.868
2.111	0.038	0.102	3.565		
job_title_Data_Engineering_Manager				1.7737	2.353
0.754	0.453	-2.919	6.467		
job_title_Data_Science_Consultant				-0.2006	1.586
-0.126	0.900	-3.363	2.962		
job_title_Data_Science_Engineer				1.1235	2.146
0.524	0.602	-3.156	5.403		
job_title_Data_Science_Manager				5.5431	1.558
3.559	0.001	2.437	8.650		
job_title_Data_Scientist				0.4854	0.844
0.575	0.567	-1.199	2.169		
job_title_Data_Specialist				1.5853	2.816
0.563	0.575	-4.032	7.203		
job_title_Director_of_Data_Engineering				2.3658	2.102
1.125	0.264	-1.827	6.558		
job_title_Director_of_Data_Science				1.5090	2.599
0.581	0.563	-3.674	6.692		
job_title_Finance_Data_Analyst				-2.6147	2.993
-0.874	0.385	-8.585	3.355		
job_title_Financial_Data_Analyst				-3.7327	3.383
-1.103	0.274	-10.479	3.014		
job_title_Head_of_Data				-4.0535	3.443

-1.177	0.243	-10.920	2.813		
job_title_Head_of_Data_Science				2.103e-11	1.87e-11
1.125	0.265	-1.63e-11	5.83e-11		
job_title_Lead_Data_Analyst				-0.5254	2.201
-0.239	0.812	-4.916	3.865		
job_title_Lead_Data_Engineer				0.9303	2.127
0.437	0.663	-3.312	5.173		
job_title_Lead_Data_Scientist				2.4796	2.972
0.834	0.407	-3.449	8.408		
job_title_ML_Engineer				1.9840	2.958
0.671	0.505	-3.916	7.885		
job_title_Machine_Learning_Engineer				1.5101	1.023
1.477	0.144	-0.529	3.550		
job_title_Machine_Learning_Infrastructure_Engineer				3.2782	2.880
1.138	0.259	-2.465	9.021		
job_title_Machine_Learning_Scientist				3.5321	3.386
1.043	0.301	-3.222	10.286		
job_title_Manager_Data_Science				1.2411	2.830
0.439	0.662	-4.404	6.886		
job_title_Marketing_Data_Analyst				2.8309	2.806
1.009	0.317	-2.766	8.428		
job_title_Principal_Data_Analyst				3.1727	2.885
1.100	0.275	-2.580	8.926		
job_title_Principal_Data_Engineer				-12.9105	4.506
-2.865	0.006	-21.898	-3.923		
job_title_Principal_Data_Scientist				3.4608	1.562
2.216	0.030	0.346	6.576		
job_title_Product_Data_Analyst				-8.5013	3.114
-2.730	0.008	-14.712	-2.290		
job_title_Research_Scientist				-0.3904	1.632
-0.239	0.812	-3.644	2.864		
job_title_Staff_Data_Scientist				-1.04e-11	1.1e-11
-0.942	0.349	-3.24e-11	1.16e-11		
employee_residence_AT				5.4057	5.513
0.981	0.330	-5.590	16.401		
employee_residence_BE				3.8739	1.555
2.491	0.015	0.773	6.975		
employee_residence_BG				4.3732	3.261
1.341	0.184	-2.130	10.877		
employee_residence_BR				-9.7715	5.431
-1.799	0.076	-20.602	1.059		
employee_residence_CA				7.6433	4.696
1.628	0.108	-1.722	17.008		
employee_residence_CL				8.527e-12	1.03e-11
0.831	0.409	-1.19e-11	2.9e-11		
employee_residence_CN				-1.6487	2.402
-0.686	0.495	-6.440	3.143		
employee_residence_CO				-0.9162	1.516

-0.604	0.548	-3.940	2.108		
employee_residence_DE				5.2784	3.439
1.535	0.129	-1.580	12.137		
employee_residence_DK				-0.6969	2.652
-0.263	0.794	-5.986	4.593		
employee_residence_ES				4.1136	2.727
1.508	0.136	-1.326	9.553		
employee_residence_FR				4.0950	2.514
1.629	0.108	-0.918	9.108		
employee_residence_GB				-4.1012	2.972
-1.380	0.172	-10.029	1.826		
employee_residence_GR				1.5076	3.174
0.475	0.636	-4.822	7.837		
employee_residence_HK				-1.585e-11	1.67e-11
-0.952	0.344	-4.91e-11	1.74e-11		
employee_residence_HR				0.5372	1.522
0.353	0.725	-2.498	3.572		
employee_residence_HU				-3.0715	4.195
-0.732	0.467	-11.439	5.296		
employee_residence_IN				-9.3334	2.367
-3.943	0.000	-14.054	-4.613		
employee_residence_IR				5.785e-12	7.21e-12
0.802	0.425	-8.6e-12	2.02e-11		
employee_residence_IT				6.1201	3.425
1.787	0.078	-0.710	12.950		
employee_residence_JE				5.2538	2.466
2.131	0.037	0.336	10.172		
employee_residence_JP				2.5501	1.000
2.550	0.013	0.555	4.545		
employee_residence_KE				-3.9494	2.099
-1.882	0.064	-8.136	0.237		
employee_residence_LU				9.024e-12	1.07e-11
0.847	0.400	-1.22e-11	3.03e-11		
employee_residence_MD				-1.2283	1.815
-0.677	0.501	-4.848	2.392		
employee_residence_MT				-1.2776	1.433
-0.892	0.376	-4.135	1.580		
employee_residence_MX				-4.2061	1.150
-3.659	0.000	-6.499	-1.913		
employee_residence_NG				-0.5572	1.041
-0.535	0.594	-2.634	1.520		
employee_residence_NL				7.1599	4.622
1.549	0.126	-2.058	16.378		
employee_residence_NZ				-1.676e-11	1.81e-11
-0.927	0.357	-5.28e-11	1.93e-11		
employee_residence_PH				0.6806	3.022
0.225	0.822	-5.347	6.708		
employee_residence_PK				-9.0381	5.540

-1.632	0.107	-20.087	2.010		
employee_residence_PL				-0.7609	2.618
-0.291	0.772	-5.983	4.461		
employee_residence_PR				5.8451	3.493
1.673	0.099	-1.122	12.813		
employee_residence_PT				2.4410	3.807
0.641	0.524	-5.152	10.034		
employee_residence_RO				-6.2619	3.747
-1.671	0.099	-13.736	1.212		
employee_residence_RS				-1.1814	4.679
-0.253	0.801	-10.513	8.150		
employee_residence_RU				16.0680	5.896
2.725	0.008	4.309	27.827		
employee_residence_SG				5.1570	1.522
3.389	0.001	2.122	8.192		
employee_residence_SI				-1.6038	1.472
-1.090	0.280	-4.539	1.331		
employee_residence_TR				-1.7678	1.120
-1.578	0.119	-4.003	0.467		
employee_residence_UA				-2.6728	1.488
-1.796	0.077	-5.640	0.295		
employee_residence_US				4.8100	1.480
3.251	0.002	1.859	7.761		
employee_residence_VN				-7.7642	4.263
-1.821	0.073	-16.266	0.737		
company_location_AS				5.6315	4.870
1.156	0.251	-4.081	15.344		
company_location_AT				0.7221	4.493
0.161	0.873	-8.239	9.684		
company_location_BE				3.8739	1.555
2.491	0.015	0.773	6.975		
company_location_BR				4.7085	6.168
0.763	0.448	-7.593	17.010		
company_location_CA				0.7494	4.269
0.176	0.861	-7.765	9.264		
company_location_CH				-0.0416	4.015
-0.010	0.992	-8.049	7.966		
company_location_CL				3.824e-17	1.44e-16
0.265	0.791	-2.49e-16	3.26e-16		
company_location_CN				3.6050	1.456
2.476	0.016	0.701	6.509		
company_location_CO				-0.9162	1.516
-0.604	0.548	-3.940	2.108		
company_location_DE				0.3665	3.602
0.102	0.919	-6.818	7.551		
company_location_DK				2.1340	1.939
1.100	0.275	-1.734	6.002		
company_location_ES				-1.3657	3.140

-0.435	0.665	-7.629	4.897		
company_location_FR				-0.8660	2.490
-0.348	0.729	-5.833	4.101		
company_location_GB				10.4463	2.778
3.760	0.000	4.905	15.988		
company_location_GR				1.1235	2.146
0.524	0.602	-3.156	5.403		
company_location_HR				0.5372	1.522
0.353	0.725	-2.498	3.572		
company_location_HU				0	0
nan	nan	0	0		
company_location_IL				5.1570	1.522
3.389	0.001	2.122	8.192		
company_location_IN				7.3119	2.343
3.120	0.003	2.638	11.985		
company_location_IR				0	0
nan	nan	0	0		
company_location_IT				0	0
nan	nan	0	0		
company_location_JP				2.5501	1.000
2.550	0.013	0.555	4.545		
company_location_KE				-3.9494	2.099
-1.882	0.064	-8.136	0.237		
company_location_LU				0	0
nan	nan	0	0		
company_location_MD				-1.2283	1.815
-0.677	0.501	-4.848	2.392		
company_location_MT				-1.2776	1.433
-0.892	0.376	-4.135	1.580		
company_location_MX				-4.2061	1.150
-3.659	0.000	-6.499	-1.913		
company_location_NG				-0.5572	1.041
-0.535	0.594	-2.634	1.520		
company_location_NL				-4.6550	5.471
-0.851	0.398	-15.567	6.257		
company_location_NZ				0	0
nan	nan	0	0		
company_location_PK				0.1987	6.375
0.031	0.975	-12.515	12.913		
company_location_PL				0.5999	2.935
0.204	0.839	-5.255	6.455		
company_location_PT				1.5683	5.002
0.314	0.755	-8.408	11.544		
company_location_RU				-4.0535	3.443
-1.177	0.243	-10.920	2.813		
company_location_SG				0	0
nan	nan	0	0		
company_location_SI				-1.6038	1.472

-1.090	0.280	-4.539	1.331		
company_location_TR				-1.7678	1.120
-1.578	0.119	-4.003	0.467		
company_location_UA				-2.6728	1.488
-1.796	0.077	-5.640	0.295		
company_location_US				2.6906	1.394
1.931	0.058	-0.089	5.470		
company_location_VN				-3.7090	4.790
-0.774	0.441	-13.263	5.845		
company_size_M				-1.5362	0.860
-1.785	0.079	-3.252	0.180		
company_size_S				-0.8028	0.962
-0.835	0.407	-2.721	1.116		

Omnibus:	24.077	Durbin-Watson:	2.067
Prob(Omnibus):	0.000	Jarque-Bera (JB):	67.549
Skew:	-0.519	Prob(JB):	2.15e-15
Kurtosis:	5.899	Cond. No.	2.14e+20

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 7.14e-29. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

```
[15]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

```

                                OLS Regression Results
=====
Dep. Variable:      cbrt_salary_in_usd    R-squared:                0.981
Model:                OLS                Adj. R-squared:          0.954
Method:              Least Squares        F-statistic:              36.26
Date:                Mon, 08 May 2023      Prob (F-statistic):       1.09e-37
Time:                00:59:09             Log-Likelihood:           -341.37
No. Observations:    171                 AIC:                     884.7
Df Residuals:        70                  BIC:                     1202.
Df Model:            100
Covariance Type:     nonrobust
=====
=====
                                coef      std err
t      P>|t|      [0.025      0.975]
-----
Intercept                                21.1044      4.810
4.388      0.000      11.512      30.697

```

salary_in_usd				9.947e-05	6.4e-06
15.532	0.000	8.67e-05	0.000		
remote_ratio				-0.0072	0.010
-0.746	0.458	-0.026	0.012		
work_year_2021e				-0.0626	0.707
-0.089	0.930	-1.473	1.348		
experience_level_EX				2.7895	2.897
0.963	0.339	-2.989	8.568		
experience_level_MI				0.8204	0.860
0.954	0.343	-0.894	2.535		
experience_level_SE				2.0696	1.071
1.932	0.057	-0.067	4.206		
employment_type_FL				-7.3516	6.302
-1.167	0.247	-19.920	5.217		
employment_type_FT				6.9583	4.491
1.549	0.126	-1.999	15.916		
employment_type_PT				7.5817	5.592
1.356	0.179	-3.571	18.734		
job_title_AI_Scientist				1.6816	2.945
0.571	0.570	-4.193	7.556		
job_title_Applied_Data_Scientist				7.3678	5.312
1.387	0.170	-3.227	17.963		
job_title_Applied_Machine_Learning_Scientist				-2.9716	3.282
-0.906	0.368	-9.517	3.574		
job_title_BI_Data_Analyst				1.5084	3.069
0.491	0.625	-4.613	7.630		
job_title_Big_Data_Architect				-0.0997	3.425
-0.029	0.977	-6.931	6.732		
job_title_Big_Data_Engineer				-1.1441	2.279
-0.502	0.617	-5.690	3.402		
job_title_Business_Data_Analyst				-5.588e-11	4.69e-11
-1.192	0.237	-1.49e-10	3.76e-11		
job_title_Cloud_Data_Engineer				-3.146e-11	3.04e-11
-1.034	0.305	-9.21e-11	2.92e-11		
job_title_Computer_Vision_Engineer				0.1228	3.424
0.036	0.971	-6.706	6.951		
job_title_Computer_Vision_Software_Engineer				1.0060	2.969
0.339	0.736	-4.916	6.928		
job_title_Data_Analyst				-0.4791	1.190
-0.403	0.688	-2.852	1.894		
job_title_Data_Analytics_Engineer				1.3717	2.883
0.476	0.636	-4.378	7.121		
job_title_Data_Analytics_Manager				2.0740	2.921
0.710	0.480	-3.751	7.900		
job_title_Data_Architect				2.9565	2.843
1.040	0.302	-2.715	8.628		
job_title_Data_Engineer				1.8334	0.868

2.111	0.038	0.102	3.565		
job_title_Data_Engineering_Manager				1.7737	2.353
0.754	0.453	-2.919	6.467		
job_title_Data_Science_Consultant				-0.2006	1.586
-0.126	0.900	-3.363	2.962		
job_title_Data_Science_Engineer				1.1235	2.146
0.524	0.602	-3.156	5.403		
job_title_Data_Science_Manager				5.5431	1.558
3.559	0.001	2.437	8.650		
job_title_Data_Scientist				0.4854	0.844
0.575	0.567	-1.199	2.169		
job_title_Data_Specialist				1.5853	2.816
0.563	0.575	-4.032	7.203		
job_title_Director_of_Data_Engineering				2.3658	2.102
1.125	0.264	-1.827	6.558		
job_title_Director_of_Data_Science				1.5090	2.599
0.581	0.563	-3.674	6.692		
job_title_Finance_Data_Analyst				-2.6147	2.993
-0.874	0.385	-8.585	3.355		
job_title_Financial_Data_Analyst				-3.7327	3.383
-1.103	0.274	-10.479	3.014		
job_title_Head_of_Data				-4.0535	3.443
-1.177	0.243	-10.920	2.813		
job_title_Head_of_Data_Science				2.103e-11	1.87e-11
1.125	0.265	-1.63e-11	5.83e-11		
job_title_Lead_Data_Analyst				-0.5254	2.201
-0.239	0.812	-4.916	3.865		
job_title_Lead_Data_Engineer				0.9303	2.127
0.437	0.663	-3.312	5.173		
job_title_Lead_Data_Scientist				2.4796	2.972
0.834	0.407	-3.449	8.408		
job_title_ML_Engineer				1.9840	2.958
0.671	0.505	-3.916	7.885		
job_title_Machine_Learning_Engineer				1.5101	1.023
1.477	0.144	-0.529	3.550		
job_title_Machine_Learning_Infrastructure_Engineer				3.2782	2.880
1.138	0.259	-2.465	9.021		
job_title_Machine_Learning_Scientist				3.5321	3.386
1.043	0.301	-3.222	10.286		
job_title_Manager_Data_Science				1.2411	2.830
0.439	0.662	-4.404	6.886		
job_title_Marketing_Data_Analyst				2.8309	2.806
1.009	0.317	-2.766	8.428		
job_title_Principal_Data_Analyst				3.1727	2.885
1.100	0.275	-2.580	8.926		
job_title_Principal_Data_Engineer				-12.9105	4.506
-2.865	0.006	-21.898	-3.923		

job_title_Principal_Data_Scientist				3.4608	1.562
2.216	0.030	0.346	6.576		
job_title_Product_Data_Analyst				-8.5013	3.114
-2.730	0.008	-14.712	-2.290		
job_title_Research_Scientist				-0.3904	1.632
-0.239	0.812	-3.644	2.864		
job_title_Staff_Data_Scientist				-1.04e-11	1.1e-11
-0.942	0.349	-3.24e-11	1.16e-11		
employee_residence_AT				5.4057	5.513
0.981	0.330	-5.590	16.401		
employee_residence_BE				3.8739	1.555
2.491	0.015	0.773	6.975		
employee_residence_BG				4.3732	3.261
1.341	0.184	-2.130	10.877		
employee_residence_BR				-9.7715	5.431
-1.799	0.076	-20.602	1.059		
employee_residence_CA				7.6433	4.696
1.628	0.108	-1.722	17.008		
employee_residence_CL				8.527e-12	1.03e-11
0.831	0.409	-1.19e-11	2.9e-11		
employee_residence_CN				-1.6487	2.402
-0.686	0.495	-6.440	3.143		
employee_residence_CO				-0.9162	1.516
-0.604	0.548	-3.940	2.108		
employee_residence_DE				5.2784	3.439
1.535	0.129	-1.580	12.137		
employee_residence_DK				-0.6969	2.652
-0.263	0.794	-5.986	4.593		
employee_residence_ES				4.1136	2.727
1.508	0.136	-1.326	9.553		
employee_residence_FR				4.0950	2.514
1.629	0.108	-0.918	9.108		
employee_residence_GB				-4.1012	2.972
-1.380	0.172	-10.029	1.826		
employee_residence_GR				1.5076	3.174
0.475	0.636	-4.822	7.837		
employee_residence_HK				-1.585e-11	1.67e-11
-0.952	0.344	-4.91e-11	1.74e-11		
employee_residence_HR				0.5372	1.522
0.353	0.725	-2.498	3.572		
employee_residence_HU				-3.0715	4.195
-0.732	0.467	-11.439	5.296		
employee_residence_IN				-9.3334	2.367
-3.943	0.000	-14.054	-4.613		
employee_residence_IR				5.785e-12	7.21e-12
0.802	0.425	-8.6e-12	2.02e-11		
employee_residence_IT				6.1201	3.425

1.787	0.078	-0.710	12.950		
employee_residence_JE				5.2538	2.466
2.131	0.037	0.336	10.172		
employee_residence_JP				2.5501	1.000
2.550	0.013	0.555	4.545		
employee_residence_KE				-3.9494	2.099
-1.882	0.064	-8.136	0.237		
employee_residence_LU				9.024e-12	1.07e-11
0.847	0.400	-1.22e-11	3.03e-11		
employee_residence_MD				-1.2283	1.815
-0.677	0.501	-4.848	2.392		
employee_residence_MT				-1.2776	1.433
-0.892	0.376	-4.135	1.580		
employee_residence_MX				-4.2061	1.150
-3.659	0.000	-6.499	-1.913		
employee_residence_NG				-0.5572	1.041
-0.535	0.594	-2.634	1.520		
employee_residence_NL				7.1599	4.622
1.549	0.126	-2.058	16.378		
employee_residence_NZ				-1.676e-11	1.81e-11
-0.927	0.357	-5.28e-11	1.93e-11		
employee_residence_PH				0.6806	3.022
0.225	0.822	-5.347	6.708		
employee_residence_PK				-9.0381	5.540
-1.632	0.107	-20.087	2.010		
employee_residence_PL				-0.7609	2.618
-0.291	0.772	-5.983	4.461		
employee_residence_PR				5.8451	3.493
1.673	0.099	-1.122	12.813		
employee_residence_PT				2.4410	3.807
0.641	0.524	-5.152	10.034		
employee_residence_RO				-6.2619	3.747
-1.671	0.099	-13.736	1.212		
employee_residence_RS				-1.1814	4.679
-0.253	0.801	-10.513	8.150		
employee_residence_RU				16.0680	5.896
2.725	0.008	4.309	27.827		
employee_residence_SG				5.1570	1.522
3.389	0.001	2.122	8.192		
employee_residence_SI				-1.6038	1.472
-1.090	0.280	-4.539	1.331		
employee_residence_TR				-1.7678	1.120
-1.578	0.119	-4.003	0.467		
employee_residence_UA				-2.6728	1.488
-1.796	0.077	-5.640	0.295		
employee_residence_US				4.8100	1.480
3.251	0.002	1.859	7.761		

employee_residence_VN				-7.7642	4.263
-1.821	0.073	-16.266	0.737		
company_location_AS				5.6315	4.870
1.156	0.251	-4.081	15.344		
company_location_AT				0.7221	4.493
0.161	0.873	-8.239	9.684		
company_location_BE				3.8739	1.555
2.491	0.015	0.773	6.975		
company_location_BR				4.7085	6.168
0.763	0.448	-7.593	17.010		
company_location_CA				0.7494	4.269
0.176	0.861	-7.765	9.264		
company_location_CH				-0.0416	4.015
-0.010	0.992	-8.049	7.966		
company_location_CL				3.824e-17	1.44e-16
0.265	0.791	-2.49e-16	3.26e-16		
company_location_CN				3.6050	1.456
2.476	0.016	0.701	6.509		
company_location_CO				-0.9162	1.516
-0.604	0.548	-3.940	2.108		
company_location_DE				0.3665	3.602
0.102	0.919	-6.818	7.551		
company_location_DK				2.1340	1.939
1.100	0.275	-1.734	6.002		
company_location_ES				-1.3657	3.140
-0.435	0.665	-7.629	4.897		
company_location_FR				-0.8660	2.490
-0.348	0.729	-5.833	4.101		
company_location_GB				10.4463	2.778
3.760	0.000	4.905	15.988		
company_location_GR				1.1235	2.146
0.524	0.602	-3.156	5.403		
company_location_HR				0.5372	1.522
0.353	0.725	-2.498	3.572		
company_location_HU				0	0
nan	nan	0	0		
company_location_IL				5.1570	1.522
3.389	0.001	2.122	8.192		
company_location_IN				7.3119	2.343
3.120	0.003	2.638	11.985		
company_location_IR				0	0
nan	nan	0	0		
company_location_IT				0	0
nan	nan	0	0		
company_location_JP				2.5501	1.000
2.550	0.013	0.555	4.545		
company_location_KE				-3.9494	2.099

-1.882	0.064	-8.136	0.237		
company_location_LU				0	0
nan	nan	0	0		
company_location_MD				-1.2283	1.815
-0.677	0.501	-4.848	2.392		
company_location_MT				-1.2776	1.433
-0.892	0.376	-4.135	1.580		
company_location_MX				-4.2061	1.150
-3.659	0.000	-6.499	-1.913		
company_location_NG				-0.5572	1.041
-0.535	0.594	-2.634	1.520		
company_location_NL				-4.6550	5.471
-0.851	0.398	-15.567	6.257		
company_location_NZ				0	0
nan	nan	0	0		
company_location_PK				0.1987	6.375
0.031	0.975	-12.515	12.913		
company_location_PL				0.5999	2.935
0.204	0.839	-5.255	6.455		
company_location_PT				1.5683	5.002
0.314	0.755	-8.408	11.544		
company_location_RU				-4.0535	3.443
-1.177	0.243	-10.920	2.813		
company_location_SG				0	0
nan	nan	0	0		
company_location_SI				-1.6038	1.472
-1.090	0.280	-4.539	1.331		
company_location_TR				-1.7678	1.120
-1.578	0.119	-4.003	0.467		
company_location_UA				-2.6728	1.488
-1.796	0.077	-5.640	0.295		
company_location_US				2.6906	1.394
1.931	0.058	-0.089	5.470		
company_location_VN				-3.7090	4.790
-0.774	0.441	-13.263	5.845		
company_size_M				-1.5362	0.860
-1.785	0.079	-3.252	0.180		
company_size_S				-0.8028	0.962
-0.835	0.407	-2.721	1.116		
=====					
Omnibus:		24.077	Durbin-Watson:		2.067
Prob(Omnibus):		0.000	Jarque-Bera (JB):		67.549
Skew:		-0.519	Prob(JB):		2.15e-15
Kurtosis:		5.899	Cond. No.		2.14e+20
=====					

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
 [2] The smallest eigenvalue is 7.14e-29. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.
 """

2.6 Define OSR2

```
[16]: # compute out-of-sample R-squared using the test set
def OSR2(model, df_train, df_test, dependent_var):
    y_test = df_test[dependent_var]
    y_pred = model.predict(df_test)
    SSE = np.sum((y_test - y_pred)**2)
    SST = np.sum((y_test - np.mean(df_train[dependent_var]))**2)
    return 1 - SSE/SST
```

```
[17]: OSR2(linreg, df_train, df_test, target_column)
```

```
[17]: 0.7394346635350122
```

2.7 Feature Selection: Use VIF to keep good features.

```
[18]: import statsmodels.api as sm
from statsmodels.stats.outliers_influence import variance_inflation_factor

def VIF(df, columns):
    values = sm.add_constant(df[columns]).values
    num_columns = len(columns) + 1
    vif = [variance_inflation_factor(values, i) for i in range(num_columns)]
    return pd.Series(vif[1:], index=columns)
```

2.8 Feature Selection: Identify and eliminate high P-Value features

```
[19]: def get_formula(features, target):
    features = [f for f in features if f != target]
    sum_features = ' + '.join(features)
    formula = ' ~ '.join([target, sum_features])
    return formula
```

```
[20]: def filter_feature(features, model, df_train, target):
    p_values = {}
    features = [f for f in model.params.index if f != 'Intercept']

    #
    for feat in features:
```

```

    p_values[feat] = model.pvalues.loc[feat]

    worst_feat = max(features, key=lambda f: p_values[f])

    print('WORST:', worst_feat, '-->', p_values[worst_feat])
    new_features = [f for f in features if f != worst_feat]

    new_model = smf.ols(formula = get_formula(new_features, target),
                        data = df_train).fit()

    return new_features, new_model, p_values[worst_feat]

```

```

[21]: features = df_train.columns
      model = linreg
      p_value = float('inf')
      models = {} # (formula, model)

      while p_value > 0.05:
          features, model, p_value = filter_feature(features, model, df_train,
          ↪target_column)
          formula = get_formula(features, target_column)

```

```

WORST: company_location_CH --> 0.991759281045693
WORST: job_title_Big_Data_Architect --> 0.9765611842494796
WORST: job_title_Data_Science_Consultant --> 0.9788993758725975
WORST: company_location_PK --> 0.9751836219950523
WORST: company_location_DE --> 0.9496630617530653
WORST: job_title_Computer_Vision_Engineer --> 0.9371771392482403
WORST: job_title_Cloud_Data_Engineer --> 0.9783733653306279
WORST: company_location_CL --> 0.98051488293352
WORST: work_year_2021e --> 0.9307019296051615
WORST: company_location_CA --> 0.9163775628226525
WORST: company_location_PL --> 0.9154613446944566
WORST: job_title_Staff_Data_Scientist --> 0.9856270115095285
WORST: company_location_AT --> 0.9138669496626517
WORST: job_title_Research_Scientist --> 0.8906949321408469
WORST: job_title_Lead_Data_Analyst --> 0.9021539749737659
WORST: employee_residence_HK --> 0.9552841025385613
WORST: job_title_Data_Analyst --> 0.8980333767796325
WORST: company_location_PT --> 0.8149601178439524
WORST: company_location_HR --> 0.8150797189764665
WORST: employee_residence_HR --> 0.8150797190256642
WORST: employee_residence_PH --> 0.9664697386596284
WORST: employee_residence_NZ --> 0.8516530309706571
WORST: employee_residence_GR --> 0.8082642133988289
WORST: employee_residence_LU --> 0.8956683389047415

```

WORST: employee_residence_IR --> 0.9906972643651115
 WORST: job_title_Business_Data_Analyst --> 0.9699199153609567
 WORST: job_title_Head_of_Data_Science --> 0.8663900157322889
 WORST: employee_residence_CL --> 0.8709368808926954
 WORST: company_location_FR --> 0.635752450341498
 WORST: company_location_ES --> 0.7197205854525648
 WORST: employee_residence_PL --> 0.6909557178948065
 WORST: company_location_HU --> 0.8195348806036717
 WORST: job_title_Big_Data_Engineer --> 0.700508935874975
 WORST: employee_residence_RS --> 0.6787171506717398
 WORST: company_location_IR --> 0.987491784671152
 WORST: company_location_IT --> 0.7093013880074468
 WORST: job_title_Lead_Data_Engineer --> 0.6564565220855878
 WORST: job_title_Computer_Vision_Software_Engineer --> 0.6686086550554402
 WORST: employee_residence_DK --> 0.6523054428769726
 WORST: job_title_Manager_Data_Science --> 0.6151040094605653
 WORST: job_title_Data_Specialist --> 0.5535398311745638
 WORST: employee_residence_NG --> 0.5415779583210076
 WORST: company_location_LU --> 0.5424528119326728
 WORST: company_location_NG --> 0.5415779580961508
 WORST: employee_residence_CO --> 0.575007133905429
 WORST: company_location_CO --> 0.5750071339086518
 WORST: job_title_Data_Analytics_Engineer --> 0.5316336596344441
 WORST: job_title_ML_Engineer --> 0.5192022493724434
 WORST: job_title_Data_Analytics_Manager --> 0.5195121681045891
 WORST: job_title_BI_Data_Analyst --> 0.5232920903626841
 WORST: employee_residence_HU --> 0.5374865851375253
 WORST: employee_residence_GB --> 0.5133020819023146
 WORST: employee_residence_MT --> 0.5055915920122626
 WORST: company_location_MT --> 0.5055915920549181
 WORST: job_title_AI_Scientist --> 0.5075837760546196
 WORST: remote_ratio --> 0.535498317312116
 WORST: employee_residence_CN --> 0.5378249025814581
 WORST: job_title_Lead_Data_Scientist --> 0.49547505799965086
 WORST: job_title_Data_Engineering_Manager --> 0.4632261450011801
 WORST: job_title_Director_of_Data_Science --> 0.5523598612457409
 WORST: job_title_Director_of_Data_Engineering --> 0.5021124501705576
 WORST: job_title_Principal_Data_Analyst --> 0.46807014702006133
 WORST: job_title_Machine_Learning_Infrastructure_Engineer -->
 0.48566268697892945
 WORST: employee_residence_SI --> 0.46050380187867246
 WORST: company_location_SI --> 0.460503801871354
 WORST: job_title_Data_Scientist --> 0.418925811063908
 WORST: job_title_Marketing_Data_Analyst --> 0.4025305400796494
 WORST: company_location_BR --> 0.41746030310367
 WORST: company_location_MD --> 0.3166879555842821
 WORST: employee_residence_MD --> 0.3166879555812153
 WORST: job_title_Data_Architect --> 0.30661237687167936

```

WORST: employee_residence_RO --> 0.31750335107150884
WORST: job_title_Machine_Learning_Engineer --> 0.21901309459394058
WORST: employee_residence_BG --> 0.23769372183538975
WORST: job_title_Machine_Learning_Scientist --> 0.2044985100407941
WORST: job_title_Principal_Data_Scientist --> 0.2395804351447887
WORST: company_location_GR --> 0.19366650816382763
WORST: job_title_Data_Science_Engineer --> 0.19366650802673335
WORST: company_location_CN --> 0.18029210215867872
WORST: employee_residence_PT --> 0.20593951540621466
WORST: company_size_M --> 0.17824332034678636
WORST: job_title_Finance_Data_Analyst --> 0.17971681822949223
WORST: company_size_S --> 0.19610909727273634
WORST: company_location_RU --> 0.15365049646812673
WORST: job_title_Head_of_Data --> 0.15365049647602907
WORST: employment_type_FL --> 0.36007373222483774
WORST: company_location_NL --> 0.14438500198718243
WORST: job_title_Data_Engineer --> 0.21137667285849612
WORST: job_title_Applied_Data_Scientist --> 0.12074026356174762
WORST: employee_residence_PR --> 0.08765391158777291
WORST: employee_residence_VN --> 0.08750858729921782
WORST: employee_residence_AT --> 0.07111522225962272
WORST: company_location_DK --> 0.07944439646037305
WORST: employee_residence_IT --> 0.09345170753000326
WORST: employee_residence_ES --> 0.1438712491672302
WORST: employee_residence_FR --> 0.18275507948126446
WORST: job_title_Applied_Machine_Learning_Scientist --> 0.0839850367891538
WORST: job_title_Financial_Data_Analyst --> 0.06962268170459038
WORST: experience_level_MI --> 0.05951294621155404
WORST: company_location_AS --> 0.06033434807966835
WORST: employee_residence_RU --> 0.042074398819232015

```

```

[22]: best_features = features + ['employee_residence_RU']
      print(best_features)
      best_linreg = smf.ols(formula=get_formula(best_features, target_column),
                           data=df_train).fit()
      print(best_linreg.summary())

```

```

['salary_in_usd', 'experience_level_EX', 'experience_level_SE',
'employment_type_FT', 'employment_type_PT', 'job_title_Data_Science_Manager',
'job_title_Principal_Data_Engineer', 'job_title_Product_Data_Analyst',
'employee_residence_BE', 'employee_residence_BR', 'employee_residence_CA',
'employee_residence_DE', 'employee_residence_IN', 'employee_residence_JE',
'employee_residence_JP', 'employee_residence_KE', 'employee_residence_MX',
'employee_residence_NL', 'employee_residence_PK', 'employee_residence_SG',
'employee_residence_TR', 'employee_residence_UA', 'employee_residence_US',
'company_location_BE', 'company_location_GB', 'company_location_IL',
'company_location_IN', 'company_location_JP', 'company_location_KE',
'company_location_MX', 'company_location_NZ', 'company_location_SG',

```

'company_location_TR', 'company_location_UA', 'company_location_US',
 'company_location_VN', 'employee_residence_RU']

OLS Regression Results

```

=====
Dep. Variable:      cbrt_salary_in_usd      R-squared:                0.966
Model:              OLS                    Adj. R-squared:          0.959
Method:             Least Squares          F-statistic:             143.4
Date:               Mon, 08 May 2023        Prob (F-statistic):      3.01e-90
Time:               00:59:28                Log-Likelihood:          -391.77
No. Observations:   171                    AIC:                     841.5
Df Residuals:       142                    BIC:                     932.7
Df Model:           28
Covariance Type:    nonrobust
=====
  
```

```

=====
                                coef      std err          t      P>|t|
-----+-----
[0.025      0.975]
-----+-----
Intercept                                23.1892         1.824        12.712      0.000
19.583      26.795
salary_in_usd                           0.0001      3.52e-06        28.838      0.000
9.46e-05      0.000
experience_level_EX                       4.2508         1.177         3.612      0.000
1.925         6.577
experience_level_SE                       2.8723         0.496         5.792      0.000
1.892         3.853
employment_type_FT                        6.7035         1.745         3.842      0.000
3.254        10.153
employment_type_PT                        5.6287         2.302         2.445      0.016
1.078        10.180
job_title_Data_Science_Manager            3.3110         1.286         2.575      0.011
0.769         5.853
job_title_Principal_Data_Engineer        -16.0164        3.072        -5.214      0.000
-22.089      -9.944
job_title_Product_Data_Analyst            -8.2223         2.746        -2.994      0.003
-13.651      -2.793
employee_residence_BE                     2.8731         1.336         2.150      0.033
0.232         5.515
employee_residence_BR                     -8.3231         1.649        -5.046      0.000
-11.584      -5.063
employee_residence_CA                     4.7788         1.408         3.394      0.001
1.996         7.562
employee_residence_DE                     3.7198         0.851         4.373      0.000
2.038         5.401
employee_residence_IN                     -9.6288         1.439        -6.692      0.000
-12.473      -6.784
employee_residence_JE                     6.3717         2.675         2.382      0.019
  
```

1.083	11.660				
employee_residence_JP		1.8942	0.711	2.664	0.009
0.489	3.299				
employee_residence_KE		-4.9128	1.333	-3.685	0.000
-7.548	-2.277				
employee_residence_MX		-4.9726	0.958	-5.189	0.000
-6.867	-3.078				
employee_residence_NL		3.9997	1.669	2.397	0.018
0.701	7.298				
employee_residence_PK		-9.9922	2.048	-4.878	0.000
-14.042	-5.943				
employee_residence_SG		3.6133	1.341	2.694	0.008
0.962	6.264				
employee_residence_TR		-2.3577	0.957	-2.465	0.015
-4.249	-0.467				
employee_residence_UA		-3.7505	1.333	-2.813	0.006
-6.386	-1.115				
employee_residence_US		3.1694	0.883	3.589	0.000
1.424	4.915				
company_location_BE		2.8731	1.336	2.150	0.033
0.232	5.515				
company_location_GB		3.3960	0.900	3.775	0.000
1.618	5.174				
company_location_IL		3.6133	1.341	2.694	0.008
0.962	6.264				
company_location_IN		5.5856	1.600	3.490	0.001
2.422	8.749				
company_location_JP		1.8942	0.711	2.664	0.009
0.489	3.299				
company_location_KE		-4.9128	1.333	-3.685	0.000
-7.548	-2.277				
company_location_MX		-4.9726	0.958	-5.189	0.000
-6.867	-3.078				
company_location_NZ		0	0	nan	nan
0	0				
company_location_SG		0	0	nan	nan
0	0				
company_location_TR		-2.3577	0.957	-2.465	0.015
-4.249	-0.467				
company_location_UA		-3.7505	1.333	-2.813	0.006
-6.386	-1.115				
company_location_US		2.1380	0.786	2.721	0.007
0.585	3.691				
company_location_VN		-14.4247	2.667	-5.408	0.000
-19.697	-9.152				
employee_residence_RU		4.3178	2.105	2.051	0.042
0.157	8.479				

=====

Omnibus:	25.929	Durbin-Watson:	1.946
Prob(Omnibus):	0.000	Jarque-Bera (JB):	43.882
Skew:	-0.772	Prob(JB):	2.96e-10
Kurtosis:	4.943	Cond. No.	5.24e+21

=====

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 1.19e-31. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

```
[23]: OSR2(best_linreg, df_train, df_test, target_column)
```

```
[23]: 0.8035179130896886
```

2.9 remove 0 coefficient and NaN p-value features

```
[24]: # best_features.remove('company_location_LU')
best_features.remove('company_location_NZ')
best_features.remove('company_location_SG')
best_linreg2 = smf.ols(formula=get_formula(best_features, target_column),
                      data=df_train).fit()
print(best_linreg2.summary())
```

OLS Regression Results

```
=====
Dep. Variable:      cbrt_salary_in_usd      R-squared:      0.966
Model:              OLS      Adj. R-squared:      0.959
Method:              Least Squares      F-statistic:      143.4
Date:                Mon, 08 May 2023      Prob (F-statistic):      3.01e-90
Time:                00:59:28      Log-Likelihood:      -391.77
No. Observations:      171      AIC:      841.5
Df Residuals:          142      BIC:      932.7
Df Model:              28
Covariance Type:      nonrobust
=====
```

```
=====
                                coef      std err          t      P>|t|
-----
[0.025      0.975]
-----
Intercept                23.1892          1.824      12.712      0.000
19.583      26.795
salary_in_usd              0.0001      3.52e-06      28.838      0.000
9.46e-05      0.000
experience_level_EX          4.2508          1.177       3.612      0.000
```


1.925	6.577				
experience_level_SE		2.8723	0.496	5.792	0.000
1.892	3.853				
employment_type_FT		6.7035	1.745	3.842	0.000
3.254	10.153				
employment_type_PT		5.6287	2.302	2.445	0.016
1.078	10.180				
job_title_Data_Science_Manager		3.3110	1.286	2.575	0.011
0.769	5.853				
job_title_Principal_Data_Engineer		-16.0164	3.072	-5.214	0.000
-22.089	-9.944				
job_title_Product_Data_Analyst		-8.2223	2.746	-2.994	0.003
-13.651	-2.793				
employee_residence_BE		2.8731	1.336	2.150	0.033
0.232	5.515				
employee_residence_BR		-8.3231	1.649	-5.046	0.000
-11.584	-5.063				
employee_residence_CA		4.7788	1.408	3.394	0.001
1.996	7.562				
employee_residence_DE		3.7198	0.851	4.373	0.000
2.038	5.401				
employee_residence_IN		-9.6288	1.439	-6.692	0.000
-12.473	-6.784				
employee_residence_JE		6.3717	2.675	2.382	0.019
1.083	11.660				
employee_residence_JP		1.8942	0.711	2.664	0.009
0.489	3.299				
employee_residence_KE		-4.9128	1.333	-3.685	0.000
-7.548	-2.277				
employee_residence_MX		-4.9726	0.958	-5.189	0.000
-6.867	-3.078				
employee_residence_NL		3.9997	1.669	2.397	0.018
0.701	7.298				
employee_residence_PK		-9.9922	2.048	-4.878	0.000
-14.042	-5.943				
employee_residence_SG		3.6133	1.341	2.694	0.008
0.962	6.264				
employee_residence_TR		-2.3577	0.957	-2.465	0.015
-4.249	-0.467				
employee_residence_UA		-3.7505	1.333	-2.813	0.006
-6.386	-1.115				
employee_residence_US		3.1694	0.883	3.589	0.000
1.424	4.915				
company_location_BE		2.8731	1.336	2.150	0.033
0.232	5.515				
company_location_GB		3.3960	0.900	3.775	0.000
1.618	5.174				
company_location_IL		3.6133	1.341	2.694	0.008

0.962	6.264				
company_location_IN		5.5856	1.600	3.490	0.001
2.422	8.749				
company_location_JP		1.8942	0.711	2.664	0.009
0.489	3.299				
company_location_KE		-4.9128	1.333	-3.685	0.000
-7.548	-2.277				
company_location_MX		-4.9726	0.958	-5.189	0.000
-6.867	-3.078				
company_location_TR		-2.3577	0.957	-2.465	0.015
-4.249	-0.467				
company_location_UA		-3.7505	1.333	-2.813	0.006
-6.386	-1.115				
company_location_US		2.1380	0.786	2.721	0.007
0.585	3.691				
company_location_VN		-14.4247	2.667	-5.408	0.000
-19.697	-9.152				
employee_residence_RU		4.3178	2.105	2.051	0.042
0.157	8.479				

```
=====
Omnibus:                25.929    Durbin-Watson:                1.946
Prob(Omnibus):           0.000    Jarque-Bera (JB):            43.882
Skew:                    -0.772    Prob(JB):                     2.96e-10
Kurtosis:                4.943    Cond. No.                     5.97e+21
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 9.16e-32. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

```
[25]: OSR2(best_linreg2, df_train, df_test, target_column)
```

```
[25]: 0.8035179130896888
```

2.10 Predict test set and compare to results

```
[26]: y_pred = best_linreg2.predict(df_test)
      # y_pred = np rint(y_pred).astype(int)
      y_pred
```

```
[26]: 201    52.284403
      14    45.656034
      12    42.383180
      226   42.813622
      140   48.904506
```

```

...
154    43.422710
18     41.290900
74     28.740143
138    47.325491
121    44.656114
Length: 74, dtype: float64

```

```
[27]: df_test
```

```

[27]:      salary_in_usd  remote_ratio  work_year_2021e  experience_level_EX  \
201          140000           100             1             0
14           103000           100             0             0
12            75966           100             1             0
226           75000            0             1             0
140          135000           100             0             0
..           ...           ...           ...           ...
154           81000            50             1             0
18           60000           100             1             0
74           28475           100             1             0
138           96357            50             1             0
121           93150            0             1             0

```

```

      experience_level_MI  experience_level_SE  employment_type_FL  \
201                    0                    1                    0
14                     1                    0                    0
12                     1                    0                    0
226                    1                    0                    0
140                    1                    0                    0
..           ...           ...           ...
154                    0                    0                    0
18                     0                    0                    0
74                     0                    0                    0
138                    0                    1                    0
121                    1                    0                    0

```

```

      employment_type_FT  employment_type_PT  job_title_AI_Scientist  ...  \
201                    1                    0                    0  ...
14                     1                    0                    0  ...
12                     1                    0                    0  ...
226                    1                    0                    0  ...
140                    1                    0                    0  ...
..           ...           ...           ...  ...
154                    1                    0                    0  ...
18                     1                    0                    0  ...
74                     1                    0                    0  ...
138                    1                    0                    0  ...

```

121	1	0	0	...
	company_location_RU	company_location_SG	company_location_SI	\
201	0	0	0	
14	0	0	0	
12	0	0	0	
226	0	0	0	
140	0	0	0	
..	
154	0	0	0	
18	0	0	0	
74	0	0	0	
138	0	0	0	
121	0	0	0	
	company_location_TR	company_location_UA	company_location_US	\
201	0	0	1	
14	0	0	1	
12	0	0	0	
226	0	0	1	
140	0	0	1	
..	
154	0	0	1	
18	0	0	1	
74	0	0	0	
138	0	0	0	
121	0	0	1	
	company_location_VN	company_size_M	company_size_S	cbirt_salary_in_usd
201	0	0	0	51.924941
14	0	0	0	46.875481
12	0	0	0	42.351918
226	0	0	0	42.171633
140	0	0	0	51.299278
..
154	0	0	1	43.267487
18	0	0	1	39.148676
74	0	1	0	30.536640
138	0	0	0	45.845258
121	0	1	0	45.330894

[74 rows x 138 columns]

2.11 predict our own entries (randomized)

read csv again with new dataframe, dropping unnecessary features

```
[28]: df_copy = pd.read_csv(r"Data Science Jobs Salaries.csv").
      ↪drop(columns=['salary_currency', 'salary_in_usd', 'salary'])
      df_copy
```

```
[28]:
```

	work_year	experience_level	employment_type	job_title \
0	2021e	EN	FT	Data Science Consultant
1	2020	SE	FT	Data Scientist
2	2021e	EX	FT	Head of Data Science
3	2021e	EX	FT	Head of Data
4	2021e	EN	FT	Machine Learning Engineer
..
240	2020	SE	FT	Data Scientist
241	2021e	MI	FT	Principal Data Scientist
242	2020	EN	FT	Data Scientist
243	2020	EN	CT	Business Data Analyst
244	2021e	SE	FT	Data Science Manager

	employee_residence	remote_ratio	company_location	company_size
0	DE	50	DE	L
1	GR	100	US	L
2	RU	0	RU	M
3	RU	50	RU	L
4	US	100	US	S
..
240	US	100	US	L
241	US	100	US	L
242	US	100	US	S
243	US	100	US	L
244	IN	50	IN	L

[245 rows x 8 columns]

get unique values of each of our features to randomize

```
[29]: possible_WY = df_copy['work_year'].unique()
      possible_ET = df_copy['employment_type'].unique()
      possible_ER = df_copy['employee_residence'].unique()
      possible_RR = df_copy['remote_ratio'].unique()
      possible_CL = df_copy['company_location'].unique()
      possible_CS = df_copy['company_size'].unique()
```

make an empty dataframe to add our imaginary data scientists to

```
[30]: random_df_copy = df_copy.copy()
      random_df_copy = random_df_copy[0:0]
      random_df_copy
```

```
[30]: Empty DataFrame
Columns: [work_year, experience_level, employment_type, job_title,
employee_residence, remote_ratio, company_location, company_size]
Index: []
```

create 200 imaginary data scientists by randomizing possible values from our raw dataset

```
[31]: for i in range(200):

    EN_JT_index = random.randint(0, len(df_copy.index)-1)

    random_WY = random.choice(possible_WY)
    random_EL = df_copy.iloc[EN_JT_index]['experience_level']
    random_ET = random.choice(possible_ET)
    random_JT = df_copy.iloc[EN_JT_index]['job_title']
    random_ER = random.choice(possible_ER)
    random_RR = random.choice(possible_RR)
    random_CL = random.choice(possible_CL)
    random_CS = random.choice(possible_CS)

    random_df_copy.loc[len(random_df_copy.index)] = [random_WY, random_EL,
    random_ET, random_JT,
    random_ER, random_RR,
    random_CL, random_CS]

random_df_copy
```

```
[31]:      work_year experience_level employment_type      job_title \
0      2021e          EN          FL      Data Scientist
1      2020          MI          CT      Data Scientist
2      2020          SE          FT      Data Engineer
3      2020          MI          CT      Data Scientist
4      2021e          EN          PT      AI Scientist
..      ...          ...          ...          ...
195    2021e          SE          FT      Data Scientist
196    2020          MI          FT      Data Scientist
197    2020          MI          FL      Data Engineer
198    2020          MI          PT  Head of Data Science
199    2020          MI          FT    Lead Data Analyst

      employee_residence  remote_ratio company_location company_size
0          PH          50          TR          M
1          BR          0          IR          S
2          SI          50          MX          M
3          FR          50          CL          M
4          HK         100          NL          S
..      ...          ...          ...          ...
```

195	PK	0	MD	L
196	PH	0	AT	M
197	GR	100	PT	S
198	BE	0	CA	S
199	NZ	0	SG	M

[200 rows x 8 columns]

create dummy variables for our features to use in making predictions

```
[32]: random_df_copy = pd.get_dummies(random_df_copy)
      random_df_copy
```

```
[32]: remote_ratio  work_year_2020  work_year_2021e  experience_level_EN \
0           50           0           1           1
1           0           1           0           0
2           50           1           0           0
3           50           1           0           0
4          100           0           1           1
..          ...           ...           ...           ...
195          0           0           1           0
196          0           1           0           0
197         100           1           0           0
198          0           1           0           0
199          0           1           0           0
```

```
experience_level_EX  experience_level_MI  experience_level_SE \
0           0           0           0
1           0           1           0
2           0           0           1
3           0           1           0
4           0           0           0
..          ...           ...           ...
195          0           0           1
196          0           1           0
197          0           1           0
198          0           1           0
199          0           1           0
```

```
employment_type_CT  employment_type_FL  employment_type_FT  ... \
0           0           1           0  ...
1           1           0           0  ...
2           0           0           1  ...
3           1           0           0  ...
4           0           0           0  ...
..          ...           ...           ...  ...
195          0           0           1  ...
```

196	0	0	1	...
197	0	1	0	...
198	0	0	0	...
199	0	0	1	...

	company_location_RU	company_location_SG	company_location_SI	\
0	0	0	0	
1	0	0	0	
2	0	0	0	
3	0	0	0	
4	0	0	0	
..	
195	0	0	0	
196	0	0	0	
197	0	0	0	
198	0	0	0	
199	0	1	0	

	company_location_TR	company_location-UA	company_location_US	\
0	1	0	0	
1	0	0	0	
2	0	0	0	
3	0	0	0	
4	0	0	0	
..	
195	0	0	0	
196	0	0	0	
197	0	0	0	
198	0	0	0	
199	0	0	0	

	company_location_VN	company_size_L	company_size_M	company_size_S
0	0	0	1	0
1	0	0	0	1
2	0	0	1	0
3	0	0	1	0
4	0	0	0	1
..
195	0	1	0	0
196	0	0	1	0
197	0	0	0	1
198	0	0	0	1
199	0	0	1	0

[200 rows x 131 columns]

rename columns to include __ in existing ones with spaces (to match the features in the model)


```
[33]: for col in random_df_copy.columns:
        random_df_copy.rename(columns={col : '_' + col.split()[0]}, inplace=True)
```

Add 0 columns for features which weren't randomly chosen from possible values. If this happens, the dummy variables for those possible categories are not in our testing set and cannot work with the model as intended.

```
[34]: for c in df.columns:
        if c not in random_df_copy.columns:
            random_df_copy[c] = 0
```

remove unneeded columns to match model perfectly

```
[35]: to_drop = ['work_year_2020', 'experience_level_EN', 'employment_type_CT',
                'job_title_3D_Computer_Vision_Researcher',
                'employee_residence_AE', 'company_location_AE', 'company_size_L',
                'log_salary_in_usd', 'sqrt_salary_in_usd', 'cbirt_salary_in_usd']
        for c in to_drop:
            if c in random_df_copy.columns:
                random_df_copy.drop(columns=[c], inplace=True)
```

print the processed, randomized dataframe of 200 data scientist

```
[36]: random_df_copy
```

```
[36]:   remote_ratio  work_year_2021e  experience_level_EX  experience_level_MI  \
0             50                1                  0                  0
1             0                 0                  0                  1
2             50                0                  0                  0
3             50                0                  0                  1
4            100                1                  0                  0
..          ...                ...                  ...                  ...
195           0                 1                  0                  0
196           0                 0                  0                  1
197          100                0                  0                  1
198           0                 0                  0                  1
199           0                 0                  0                  1
```

```
   experience_level_SE  employment_type_FL  employment_type_FT  \
0                   0                    1                    0
1                   0                    0                    0
2                   1                    0                    1
3                   0                    0                    0
4                   0                    0                    0
..          ...                ...                  ...
195                  1                    0                    1
196                  0                    0                    1
197                  0                    1                    0
```

198	0	0	0
199	0	0	1

	employment_type_PT	job_title_AI_Scientist	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	1	1	
..	
195	0	0	
196	0	0	
197	0	0	
198	1	0	
199	0	0	

	job_title_Applied_Data_Scientist	...	job_title_Cloud_Data_Engineer	\
0	0	...	0	
1	0	...	0	
2	0	...	0	
3	0	...	0	
4	0	...	0	
..	
195	0	...	0	
196	0	...	0	
197	0	...	0	
198	0	...	0	
199	0	...	0	

	job_title_Data_Architect	job_title_Data_Science_Engineer	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	
..	
195	0	0	
196	0	0	
197	0	0	
198	0	0	
199	0	0	

	job_title_Data_Specialist	job_title_Director_of_Data_Engineering	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	

4	0	0
..
195	0	0
196	0	0
197	0	0
198	0	0
199	0	0

	job_title_Finance_Data_Analyst	job_title_Financial_Data_Analyst	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	
..	
195	0	0	
196	0	0	
197	0	0	
198	0	0	
199	0	0	

	job_title_Machine_Learning_Infrastructure_Engineer	\
0	0	
1	0	
2	0	
3	0	
4	0	
..	...	
195	0	
196	0	
197	0	
198	0	
199	0	

	job_title_Marketing_Data_Analyst	job_title_Principal_Data_Analyst
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0
..
195	0	0
196	0	0
197	0	0
198	0	0
199	0	0


[200 rows x 137 columns]

make predictions for our 200 random data scientist's salaries

```
[37]: y_pred_rand = best_linreg2.predict(random_df_copy)
      # y_pred_rand = np rint(y_pred_rand).astype(int)
      y_pred_rand
```

```
[37]: 0      20.831496
      1      14.866011
      2      27.792390
      3      23.189158
      4      28.817815
      ...
      195     22.772765
      196     29.892684
      197     23.189158
      198     31.690948
      199     29.892684
      Length: 200, dtype: float64
```

```
[38]: random_df_copy['pred_cbrt_salary_in_usd'] = y_pred_rand
      sns.displot(data=random_df_copy, x = 'pred_cbrt_salary_in_usd')
      plt.show();
```



```
[39]: # random_df_copy['pred_salary'] = y_pred_rand
negative_predictions = random_df_copy[random_df_copy['pred_cbrt_salary_in_usd'] <
↪ 0]
display(negative_predictions)
```

Empty DataFrame

```

Columns: [remote_ratio, work_year_2021e, experience_level_EX,
↳experience_level_MI, experience_level_SE, employment_type_FL,
↳employment_type_FT, employment_type_PT, job_title_AI_Scientist,
↳job_title_Applied_Data_Scientist, job_title_Big_Data_Architect,
↳job_title_Big_Data_Engineer, job_title_Business_Data_Analyst,
↳job_title_Computer_Vision_Engineer,
↳job_title_Computer_Vision_Software_Engineer, job_title_Data_Analyst,
↳job_title_Data_Analytics_Engineer, job_title_Data_Analytics_Manager,
↳job_title_Data_Engineer, job_title_Data_Engineering_Manager,
↳job_title_Data_Science_Consultant, job_title_Data_Science_Manager,
↳job_title_Data_Scientist, job_title_Director_of_Data_Science,
↳job_title_Head_of_Data, job_title_Head_of_Data_Science,
↳job_title_Lead_Data_Analyst, job_title_Lead_Data_Engineer,
↳job_title_Lead_Data_Scientist, job_title_ML_Engineer,
↳job_title_Machine_Learning_Engineer, job_title_Machine_Learning_Scientist,
↳job_title_Manager_Data_Science, job_title_Principal_Data_Engineer,
↳job_title_Principal_Data_Scientist, job_title_Product_Data_Analyst,
↳job_title_Research_Scientist, job_title_Staff_Data_Scientist,
↳employee_residence_AT, employee_residence_BE, employee_residence_BG,
↳employee_residence_BR, employee_residence_CA, employee_residence_CL,
↳employee_residence_CN, employee_residence_CO, employee_residence_DE,
↳employee_residence_DK, employee_residence_ES, employee_residence_FR,
↳employee_residence_GB, employee_residence_GR, employee_residence_HK,
↳employee_residence_HR, employee_residence_HU, employee_residence_IN,
↳employee_residence_IR, employee_residence_IT, employee_residence_JE,
↳employee_residence_JP, employee_residence_KE, employee_residence_LU,
↳employee_residence_MD, employee_residence_MT, employee_residence_MX,
↳employee_residence_NG, employee_residence_NL, employee_residence_NZ,
↳employee_residence_PH, employee_residence_PK, employee_residence_PL,
↳employee_residence_PR, employee_residence_PT, employee_residence_RO,
↳employee_residence_RS, employee_residence_RU, employee_residence_SG,
↳employee_residence_SI, employee_residence_TR, employee_residence_UA,
↳employee_residence_US, employee_residence_VN, company_location_AS,
↳company_location_AT, company_location_BE, company_location_BR,
↳company_location_CA, company_location_CH, company_location_CL,
↳company_location_CN, company_location_CO, company_location_DE,
↳company_location_DK, company_location_ES, company_location_FR,
↳company_location_GB, company_location_GR, company_location_HR,
↳company_location_HU, company_location_IL, ...]

Index: []

```

```
[0 rows x 138 columns]
```

```
[40]: #plt.plot(negative_predictions);
```

```
[41]: x_min = np.min(y_pred)
x_max = np.max(y_pred)
```


```

y_min = np.min(df_test_y_actual)
y_max = np.max(df_test_y_actual)

x = np.arange(x_min, x_max + 1)
y = np.arange(y_min, y_max + 1)

plt.scatter(x=y_pred, y=df_test_y_actual, s=15)
plt.plot(x, y, color='red') # y = x
plt.xlabel('y_predicted')
plt.ylabel('y_actual')
plt.title('y_predicted vs y_actual')
plt.show();

```



```
[42]: residuals = df_test_y_actual - y_pred
```

```

x_min = np.min(y_pred)
x_max = np.max(y_pred)
y_min = np.min(residuals)
y_max = np.max(residuals)

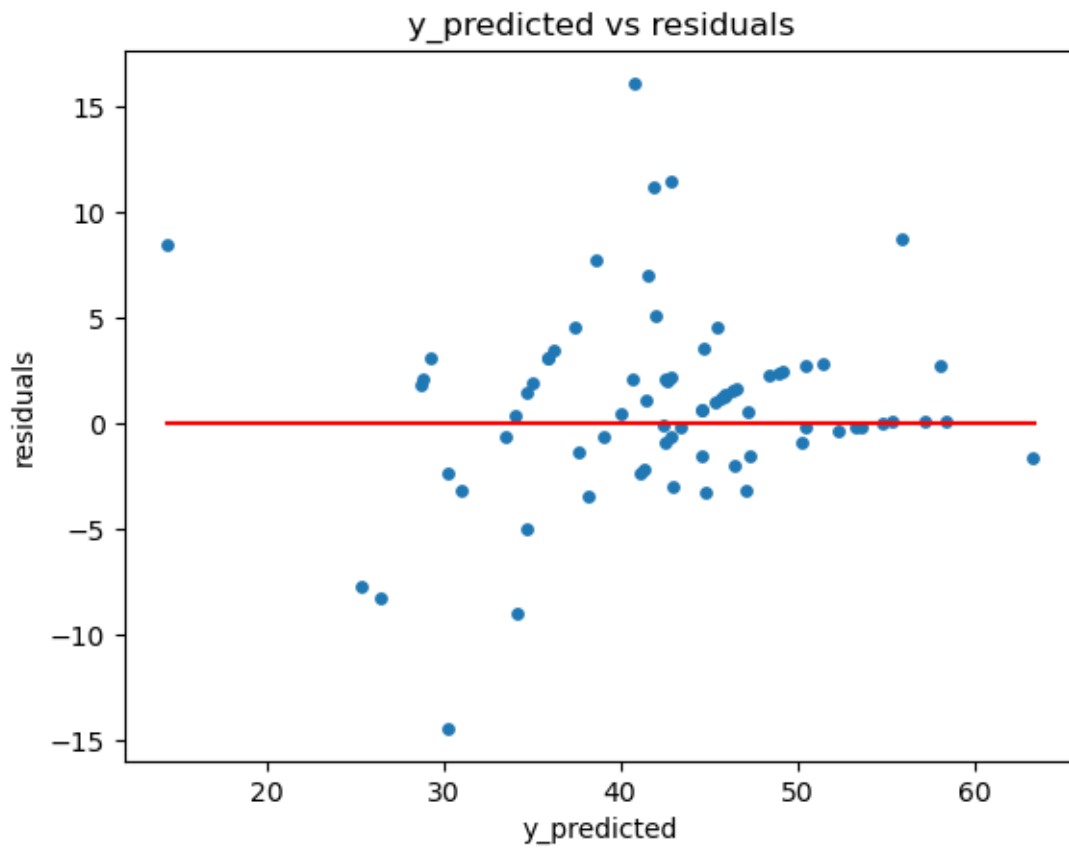
```

```

x = np.arange(x_min, x_max + 1)
y = np.zeros(len(x))

plt.scatter(x=y_pred, y=residuals, s=15)
plt.plot(x, y, color='red') #  $y = x$ 
plt.xlabel('y_predicted')
plt.ylabel('residuals')
plt.title('y_predicted vs residuals')
plt.show();

```



```
[43]: error = df_test_y_actual - y_pred
```

```
[ ]:
```